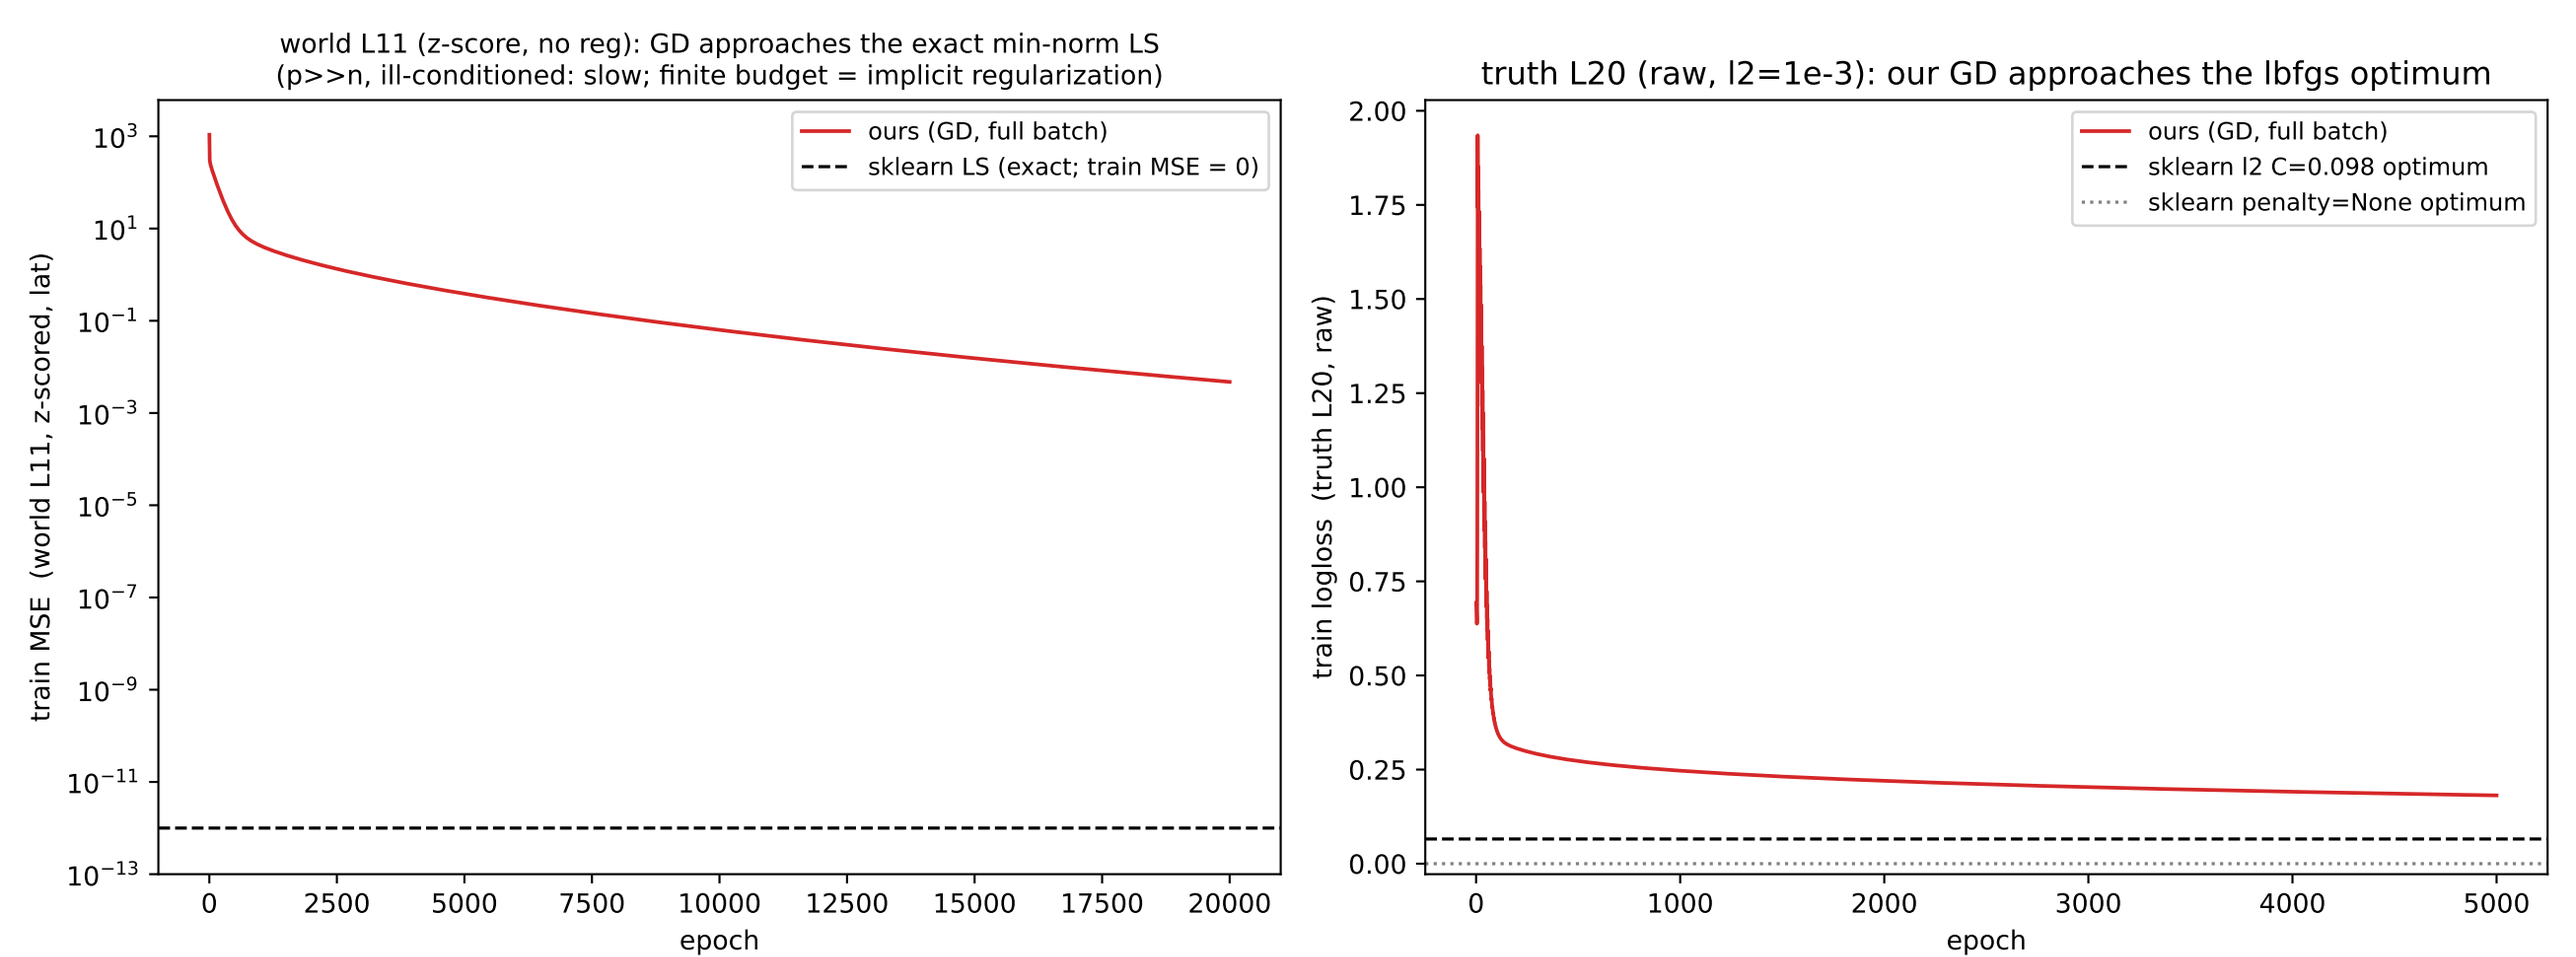
world L11 (z-score, no reg): GD approaches the exact min-norm LS
(p>>n, ill-conditioned: slow; finite budget = implicit regularization)
ours (GD, full batch)
sklearn LS (exact; train MSE = 0)
train MSE (world L11, z-scored, lat)
epoch
truth L20 (raw, l2=1e-3): our GD approaches the lbfgs optimum
ours (GD, full batch)
sklearn l2 C=0.098 optimum
sklearn penalty=None optimum
train logloss (truth L20, raw)
epoch